CS187 Lab 4-5: Sequence-to-sequence models with attention

September 7, 2024

```
[]: # Please do not change this cell because some hidden tests might depend on it.
     import os
     # Otter grader does not handle ! commands well, so we define and use our
     # own function to execute shell commands.
     def shell(commands, warn=True):
         """Executes the string `commands` as a sequence of shell commands.
            Prints the result to stdout and returns the exit status.
            Provides a printed warning on non-zero exit status unless `warn`
           flag is unset.
         file = os.popen(commands)
         print (file.read().rstrip('\n'))
         exit_status = file.close()
         if warn and exit_status != None:
             print(f"Completed with errors. Exit status: {exit_status}\n")
         return exit_status
     shell("""
     ls requirements.txt >/dev/null 2>&1
     if [ ! $? = 0 ]; then
     rm -rf .tmp
     git clone https://github.com/cs187-2021/lab4-5.git .tmp
     mv .tmp/tests ./
     mv .tmp/requirements.txt ./
     rm -rf .tmp
     fi
     pip install -q -r requirements.txt
```

```
[]: # Initialize Otter
import otter
grader = otter.Notebook()
```

In lab 4-4, you built a sequence-to-sequence model in its most basic form and applied it to the task of words-to-numbers conversion. That model first encodes the source sequence into a fixed-size vector (encoder final states), and then decodes based on that vector. Since the only way information from the source side can flow to the target side is through this fixed-size vector, it presents a bottleneck in the encoder-decoder model: no matter how long the source sentence is, it must always be compressed into this fixed-size vector.

An attention mechanism (proposed in this seminal paper) offers a workaround by providing the decoder a dynamic view of the source-side as the decoding proceeds. Instead of compressing the source sequence into a fixed-size vector, we preserve the "resolution" and encode the source sequence into a set of vectors (usually with the same size as the source sequence) which is sometimes called a memory bank. When predicting each word, the decoder "attends to" this memory bank and assigns a weight to each vector in the set, and the weighted sum of those vectors will be used to make a prediction. Hopefully, the decoder will assign higher weights to more relevant source words when predicting a target word, which we'll test in this lab.

New bits of Pytorch used in this lab, and which you may find useful include:

- torch.transpose: swaps two dimensions of a tensor.
- torch.reshape: reshapes a tensor.
- torch.bmm: Performs batched matrix multiplication.
- torch.nn.utils.rnn.pack_padded_sequence (imported as pack): Handles paddings. A more detailed explanation can be found here.
- torch.nn.utils.rnn.pad_packed_sequence (imported as unpack): Handles paddings.
- torch.masked_fill: Fills tensor elements with a value in spots where mask is True.
- torch.softmax: Computes softmax.
- torch.repeat: Repeats a tensor along the specified dimensions.
- torch.triu: Returns the upper triangular part of a matrix.

Preparation - Loading data

We use the same data as in lab 4-4.

```
import copy
import csv
import math
import matplotlib
import matplotlib.pyplot as plt
import os
import wget

import torch
import torch.nn as nn

from datasets import load_dataset

from tokenizers import Tokenizer
from tokenizers.pre_tokenizers import WhitespaceSplit
from tokenizers.processors import TemplateProcessing
```

```
from tokenizers import normalizers
from tokenizers.models import WordLevel
from tokenizers.trainers import WordLevelTrainer
from transformers import PreTrainedTokenizerFast

from tqdm import tqdm

from torch.nn.utils.rnn import pack_padded_sequence as pack
from torch.nn.utils.rnn import pad_packed_sequence as unpack
```

```
[]: # Specify matplotlib configuration
%matplotlib inline
plt.style.use('tableau-colorblind10')

# GPU check, make sure to use GPU where available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)
```

```
[]: # Download data
     def download_if_needed(source, dest, filename):
         os.makedirs(dest, exist_ok=True) # ensure destination
         os.path.exists(f''./{dest}{filename}") or wget.download(source + filename,__
      →out=dest)
     local_dir = "data/"
     remote_dir = "https://github.com/nlp-course/data/raw/master/Words2Num/"
     os.makedirs(local_dir, exist_ok=True)
     for filename in [
         "train.src".
         "train.tgt",
         "dev.src",
         "dev.tgt",
         "test.src",
         "test.tgt",
         download_if_needed(remote_dir, local_dir, filename)
```

As in lab 4-4, we process the dataset by extracting the sequences and their corresponding labels and save it in CSV format. Then, we load the data from the CSV files, train the tokenizers, prepend

<

```
[]: # Process data
for split in ['train', 'dev', 'test']:
    src_in_file = f'{local_dir}{split}.src'
    tgt_in_file = f'{local_dir}{split}.tgt'
    out_file = f'{local_dir}{split}.csv'
```

```
with open(src_in_file, 'r') as f_src_in, open(tgt_in_file, 'r') as f_tgt_in:
                 with open(out_file, 'w') as f_out:
                          src, tgt= [], []
                         writer = csv.writer(f_out)
                         writer.writerow(('src','tgt'))
                         for src_line, tgt_line in zip(f_src_in, f_tgt_in):
                                  writer.writerow((src_line.strip(), tgt_line.strip()))
dataset = load_dataset('csv', data_files={'train':f'{local_dir}train.csv', \
                                                                                           'val': f'{local dir}dev.csv', \
                                                                                           'test': f'{local_dir}test.csv'})
train_data = dataset['train']
val_data = dataset['val']
test_data = dataset['test']
unk_token = '[UNK]'
pad_token = '[PAD]'
bos_token = '<bos>'
eos_token = '<eos>'
src_tokenizer = Tokenizer(WordLevel(unk_token=unk_token))
src_tokenizer.pre_tokenizer = WhitespaceSplit()
src_trainer = WordLevelTrainer(special_tokens=[pad_token, unk_token])
src_tokenizer.train_from_iterator(train_data['src'], trainer=src_trainer)
tgt_tokenizer = Tokenizer(WordLevel(unk_token=unk_token))
tgt_tokenizer.pre_tokenizer = WhitespaceSplit()
tgt_trainer = WordLevelTrainer(special_tokens=[pad_token, unk_token, bos_token, bos_toke
  ⇔eos_token])
tgt_tokenizer.train_from_iterator(train_data['tgt'], trainer=tgt_trainer)
tgt_tokenizer.post_processor = TemplateProcessing(single=f"{bos_token} $A_\_
  -token_to_id(bos_token)), (eos_token, tgt_tokenizer.token_to_id(eos_token))])
hf_src_tokenizer = PreTrainedTokenizerFast(tokenizer_object=src_tokenizer,_
   →pad_token=pad_token, unk_token=unk_token)
hf tgt tokenizer = PreTrainedTokenizerFast(tokenizer object=tgt tokenizer,
   →pad_token=pad_token, unk_token=unk_token, bos_token=bos_token, unk_token
  ⇔eos token=eos token)
def encode(example):
        example['src_ids'] = hf_src_tokenizer(example['src']).input_ids
```

```
example['tgt_ids'] = hf_tgt_tokenizer(example['tgt']).input_ids
    return example

train_data = train_data.map(encode)
val_data = val_data.map(encode)

test_data = test_data.map(encode)

# Compute size of vocabulary
src_vocab = src_tokenizer.get_vocab()
tgt_vocab = tgt_tokenizer.get_vocab()

print(f"Size of src vocab: {len(src_vocab)}")
print(f"Size of tgt vocab: {len(tgt_vocab)}")
print(f"Index for src padding: {src_vocab[pad_token]}")
print(f"Index for start of sequence token: {tgt_vocab[bos_token]}")
print(f"Index for end of sequence token: {tgt_vocab[eos_token]}")
```

To load data in batched tensors, we use torch.utils.data.DataLoader for data splits, which enables us to iterate over the dataset under a given BATCH_SIZE. For the test set, we use a batch size of 1, to make the decoding implementation easier.

```
[ ]: BATCH_SIZE = 32
                        # batch size for training and validation
     TEST_BATCH_SIZE = 1 # batch size for test; we use 1 to make implementation_
      \rightarrow easier
     # Defines how to batch a list of examples together
     def collate_fn(examples):
         batch = \{\}
         bsz = len(examples)
         src_ids, tgt_ids = [], []
         for example in examples:
             src_ids.append(example['src_ids'])
             tgt_ids.append(example['tgt_ids'])
         src_len = torch.LongTensor([len(word_ids) for word_ids in src_ids]).
      →to(device)
         src_max_length = max(src_len)
         tgt_max_length = max([len(word_ids) for word_ids in tgt_ids])
         src_batch = torch.zeros(bsz, src_max_length).long().
      →fill_(src_vocab[pad_token]).to(device)
         tgt_batch = torch.zeros(bsz, tgt_max_length).long().

¬fill_(tgt_vocab[pad_token]).to(device)

         for b in range(bsz):
             src_batch[b][:len(src_ids[b])] = torch.LongTensor(src_ids[b]).to(device)
             tgt_batch[b][:len(tgt_ids[b])] = torch.LongTensor(tgt_ids[b]).to(device)
```

```
batch['src_lengths'] = src_len
    batch['src_ids'] = src_batch
    batch['tgt_ids'] = tgt_batch
    return batch
train_iter = torch.utils.data.DataLoader(train_data,
                                          batch_size=BATCH_SIZE,
                                          shuffle=True,
                                          collate_fn=collate_fn)
val iter = torch.utils.data.DataLoader(val data,
                                        batch_size=BATCH_SIZE,
                                        shuffle=False,
                                        collate_fn=collate_fn)
test_iter = torch.utils.data.DataLoader(test_data,
                                         batch_size=TEST_BATCH_SIZE,
                                         shuffle=False,
                                         collate_fn=collate_fn)
```

Let's take a look at a batch from these iterators.

```
batch = next(iter(train_iter))
src_ids = batch['src_ids']
src_example = src_ids[2]
print (f"Size of src batch: {src_ids.size()}")
print (f"Third src sentence in batch: {src_example}")
print (f"Length of the third src sentence in batch: {len(src_example)}")
print (f"Converted back to string: {hf_src_tokenizer.decode(src_example)}")

tgt_ids = batch['tgt_ids']
tgt_example = tgt_ids[2]
print (f"Size of tgt batch: {tgt_ids.size()}")
print (f"Third tgt sentence in batch: {tgt_example}")
print (f"Converted back to string: {hf_tgt_tokenizer.decode(tgt_example)}")
```

1 The attention mechanism

Attention works by querying a (dynamically sized) set of keys associated with values. As usual, the query, keys, and values are represented as vectors. The query process provides a score that specifies how much each key should be attended to. The attention can then be summarized by taking an average of the values weighted by the attention score of the corresponding keys. This context vector can then be used as another input to other processes.

More formally, let's suppose we have a query vector $\mathbf{q} \in \mathbb{R}^D$, a set of S key-value pairs $\{(\mathbf{k}_i, \mathbf{v}_i) \in \mathbb{R}^D \times \mathbb{R}^D : i \in \{1, 2, \cdots, S\}\}$, where D is the hidden size. What we want to do through the attention mechanism is to use the query to attend to the keys, and summarize those values associated with the "relevant" keys into a fixed-size context vector $\mathbf{c} \in \mathbb{R}^D$. Note that this is different from directly

compressing the key-value pairs into a fixed-size vector, since depending on the query, we might end up with different context vectors.

To determine the score for a given query and key, it is standard to use a measure of similarity between the query and key. You've seen such similarity measures before, in labs 1-1 and 1-2. A good choice is simply the normalized dot product between query and key. We'll thus take the attention score for query \mathbf{q} and key \mathbf{k}_i to be

$$a_i = \frac{\exp(\mathbf{q} \cdot \mathbf{k}_i)}{Z},$$

where \cdot denotes the dot product (inner product) and exp is exponentiation which ensures that all scores are nonnegative, and

$$Z = \sum_{i=1}^{S} \exp(\mathbf{q} \cdot \mathbf{k}_i)$$

is the normalizer to guarantee the scores all sum to one. (There are multiple ways of parameterizing the attention function, but the form we present here is the most popular one.) You might have noticed that the operation above is essentially a softmax over $\mathbf{q} \cdot \mathbf{k}$.

The attention scores **a** lie on a *simplex* (meaning $a_i \ge 0$ and $\sum_i a_i = 1$), which lends it some interpretability: the closer a_i is to 1, the more "relevant" a key k_i (and hence its value v_i) is to the given query. We will observe this later in the lab: When we are about to predict the target word "3", a_i is close to 1 for the source word x_i = "three".

To compute the context vector \mathbf{c} , we take the weighted sum of values using the corresponding attention scores as weights:

$$\mathbf{c} = \sum_{i=1}^{S} a_i \mathbf{v}_i$$

The closer a_i is to 1, the higher the weight \mathbf{v}_i receives.

Question: In the extreme, if there exists i for which a_i is 1, then what will the value of \mathbf{c} be? Type your answer here, replacing this text.

In practice, instead of computing the context vector once for each query, we want to batch computations for different queries together for parallel processing on GPUs. This will become especially useful for the transformer implementation. We use a matrix $Q \in \mathbb{R}^{T \times D}$ to store T queries, a matrix $K \in \mathbb{R}^{S \times D}$ to store S keys, and a matrix $V \in \mathbb{R}^{S \times D}$ to store the corresponding values. Then we can write down how we compute the attention scores $A \in \mathbb{R}^{T \times S}$ in a matrix form:

$$A = \operatorname{softmax}(QK^\top, \dim = -1),$$

Question: What is the shape of A? What does A_{ij} represent?

Type your answer here, replacing this text.

To get the context matrix $C \in \mathbb{R}^{T \times D}$:

Your first job is to implement this calculation by finishing the attention function below, which takes the Q, K, and V matrices and returns the A and C matrices. Note that for these matrices, there is one additional dimension for the batching, so instead of $Q \in \mathbb{R}^{T \times D}$, $K, V \in \mathbb{R}^{S \times D}$, $A \in \mathbb{R}^{T \times S}$, $C \in \mathbb{R}^{T \times D}$, we have $Q \in \mathbb{R}^{B \times T \times D}$, $K, V \in \mathbb{R}^{B \times S \times D}$, $K, V \in \mathbb{R}^{B \times T \times S}$, $K, V \in \mathbb{R}^{B \times T \times D}$, where $K, V \in \mathbb{R}^{B \times T \times S}$ is the batch size. In addition, the function below also takes an argument mask of size $K, V \in \mathbb{R}^{B \times T \times S}$ to mark where attentions are disallowed. This is useful not only in disallowing attending to padding symbols, but also in implementing the transformer model which we'll see later in this lab.

Hint: You might find torch.bmm helpful for batched matrix multiplications. You might need to transpose and reshape tensors to be able to use this function.

Hint: As mentioned in the beginning of the lab, you might also find torch.transpose, torch.reshape, torch.masked_fill, and torch.softmax useful.

Hint: A simple trick for masking an attention score is to set it to negative infinity before normalization.

```
[]: #TODO - finish implementing this function.
     def attention(batched_Q, batched_K, batched_V, mask=None):
       Performs the attention operation and returns the attention matrix
       `batched_A` and the context matrix `batched_C` using queries
       `batched_Q`, keys `batched_K`, and values `batched_V`.
       Arguments:
           batched_Q: (bsz, q_len, D)
           batched_K: (bsz, k_len, D)
           batched V: (bsz, k len, D)
           mask: (bsz, q_len, k_len). An optional boolean mask *disallowing*
                 attentions where the mask value is *`False`*.
       Returns:
           batched_A: the normalized attention scores (bsz, q_len, k_len)
           batched_C: a tensor of size (bsz, q_len, D).
       # Check sizes
       D = batched_Q.size(-1)
       bsz = batched_Q.size(0)
       q_len = batched_Q.size(1)
       k_len = batched_K.size(1)
       assert batched_K.size(-1) == D and batched_V.size(-1) == D
       assert batched_K.size(0) == bsz and batched_V.size(0) == bsz
       assert batched_V.size(1) == k_len
       if mask is not None:
         assert mask.size() == torch.Size([bsz, q_len, k_len])
       batched A = \dots
```

```
[]: grader.check("attention")
```

1.1 Neural encoder-decoder models with attention

Now we can add an attention mechanism to our encoder-decoder model. As in lab 4-4, we use a bidirectional LSTM as the encoder, and a unidirectional LSTM as the decoder, and initialize the decoder state with the encoder final state. However, instead of directly projecting the decoder hidden state to logits, we use it as a query vector and attend to all encoder outputs (used as both keys and values), and then concatanate the resulting context vector with the query vector, and project to logits. In addition, we add the context vector to the word embedding at the next time step, so that the LSTM can be aware of the previous attention results.

In the above illustration, at the first time step, we use q_1 to denote the decoder output. Instead of directly projecting that to logits as in lab 4-4, we use q_1 as the query vector, and use it to attend to the memory bank (which is the set of encoder outputs) and get the context vector c_1 . We concatenate c_1 with q_1 , and project the result to the vocabulary size to get logits. At the next step, we first embed y_1 into embeddings, and then $\mathbf{add}\ c_1$ to it (via componentwise addition) and use the sum as the decoder input. This process continues until an end-of-sequence is produced.

You'll need to implement forward_encoder and forward_decoder_incrementally in the code below. The forward_encoder function will return a "memory bank" in addition to the final states. The "memory bank" is simply the encoder outputs at all time steps, which is the first returned value of torch.nn.LSTM.

The forward_decoder_incrementally function forwards the LSTM cell for a single time step. It takes the initial decoder state, the memory bank, and the input word at the current time step and returns logits for this time step. In addition, it needs to return the context vector and the updated decoder state, which will be used for the next time step. Note that here you need to consider batch sizes greater than 1, as this function is used in forward_decoder, which is used during training.

In summary, the steps in decoding are:

- 1. Map the target words to word embeddings. Add the context vector from the previous time step if any. Use the result as the input to the decoder.
- 2. Forward the decoder RNN for one time step. Use the decoder output as query, the memory bank as **both keys and values**, and compute the context vector through the attention mechanism. Since we don't want to attend to padding symbols at the source side, we also need to pass in a proper mask to the attention function.
- 3. Concatenate the context vector with the decoder output, and project the concatenation to vocabulary size as (unnormalized) logits. Normalize them using torch.log_softmax if normalize is True.

4. Update the decoder hidden state and the context vector, which will be used in the next time step.

Before proceeding, let's consider a simple question: in lab 4-4, we tried to avoid for loops, but if you read the code of forward_decoder in this lab, you might notice a for loop. Is this unavoidable?

Question: Recall that in the forward_decoder function in lab 4-4 we didn't use any for loops but instead used a single call to self.decoder_rnn. Why do we need a for loop in the function forward_decoder below? Is it possible to get rid of the for loop to make the code more efficient?

Type your answer here, replacing this text.

Now let's implement forward_encoder and forward_decoder_incrementally.

Hint on using pack: if you use pack to handle paddings and pass the result as encoder inputs, you need to use unpack and extract the first returned value as the memory bank. An example can be found here, but note that our input is already the padded sequences, and that we set batch_first to False. Hint on ignoring source-side paddings in the attention mechanism: what mask should we pass into the attention function??

```
[]: #TODO - implement `forward_encoder` and `forward_decoder_incrementally`.
     class AttnEncoderDecoder(nn.Module):
       def __init__(self, hf_src_tokenizer, hf_tgt_tokenizer, hidden_size=64,_
      →lavers=3):
         11 11 11
         Initializer. Creates network modules and loss function.
         Arguments:
             hf_src_tokenizer: hf src tokenizer
             hf_tqt_tokenizer: hf tqt tokenizer
             hidden_size: hidden layer size of both encoder and decoder
             layers: number of layers of both encoder and decoder
         super().__init__()
         self.hf_src_tokenizer = hf_src_tokenizer
         self.hf_tgt_tokenizer = hf_tgt_tokenizer
         # Keep the vocabulary sizes available
         self.V_src = len(self.hf_src_tokenizer)
         self.V_tgt = len(self.hf_tgt_tokenizer)
         # Get special word ids
         self.padding_id_src = self.hf_src_tokenizer.pad_token_id
         self.padding_id_tgt = self.hf_tgt_tokenizer.pad_token_id
         self.bos_id = self.hf_tgt_tokenizer.bos_token_id
         self.eos_id = self.hf_tgt_tokenizer.eos_token_id
         # Keep hyper-parameters available
```

```
self.embedding_size = hidden_size
  self.hidden_size = hidden_size
  self.layers = layers
  # Create essential modules
  self.word_embeddings_src = nn.Embedding(self.V_src, self.embedding_size)
  self.word_embeddings_tgt = nn.Embedding(self.V_tgt, self.embedding_size)
  # RNN cells
  self.encoder rnn = nn.LSTM(
    input size = self.embedding size,
    hidden_size = hidden_size // 2, # to match decoder hidden size
    num layers
                 = layers,
    batch_first=True,
    bidirectional = True
                                    # bidirectional encoder
  self.decoder_rnn = nn.LSTM(
    input_size = self.embedding_size,
    hidden_size = hidden_size,
    num_layers
                = layers,
    batch_first=True,
    )
  # Final projection layer
  self.hidden2output = nn.Linear(2*hidden_size, self.V_tgt) # project the
⇔concatenation to logits
  # Create loss function
  self.loss_function = nn.CrossEntropyLoss(reduction='sum',
                                          ignore_index=self.padding_id_tgt)
def forward_encoder(self, src, src_lengths):
  Encodes source words `src`.
  Arguments:
      src: src batch of size (bsz, max_src_len)
      src_lengths: src lengths of size (bsz)
  Returns:
      memory_bank: a tensor of size (bsz, src_len, hidden_size)
      (final_state, context): `final_state` is a tuple (h, c) where h/c is of \sqcup
\hookrightarrow size
                              (layers, bsz, hidden_size), and `context` is\Box
→ `None`.
  11 11 11
  #TODO
```

```
memory_bank = ...
  final_state = ...
   context = None
  return memory_bank, (final_state, context)
def forward_decoder(self, encoder_final_state, tgt_in, memory_bank, src_mask):
  Decodes based on encoder final state, memory bank, src_mask, and ground_{\sqcup}
\hookrightarrow truth
   target words.
  Arguments:
       encoder\_final\_state: (final\_state, None) where final\_state is the \sqcup
\hookrightarrow encoder
                              final state used to initialize decoder. None is the
                              initial context (there's no previous context at the
                              first step).
       tgt_in: a tensor of size (bsz, tgt_len)
       memory_bank: a tensor of size (bsz, src_len, hidden_size), encoder_\_
\hookrightarrow outputs
                     at every position
       src_mask: a tensor of size (bsz, src_len): a boolean tensor, `False`⊔
\hookrightarrow where
                  src is padding (we disallow decoder to attend to those ⊔
\hookrightarrow places).
  Returns:
       Logits of size (bsz, tgt_len, V_tgt) (before the softmax operation)
  max_tgt_length = tgt_in.size(1)
  # Initialize decoder state, note that it's a tuple (state, context) here
  decoder_states = encoder_final_state
  all logits = []
  for i in range(max_tgt_length):
     logits, decoder_states, attn = \
       self.forward_decoder_incrementally(decoder_states,
                                             tgt_in[:, i],
                                             memory_bank,
                                             src_mask,
                                             normalize=False)
     all_logits.append(logits)
                                             # list of bsz, vocab_tqt
  all_logits = torch.stack(all_logits, 1) # bsz, tgt_len, vocab_tgt
  return all_logits
def forward(self, src, src_lengths, tgt_in):
  Performs forward computation, returns logits.
```

```
Arguments:
       src: src batch of size (bsz, max_src len)
       src_lengths: src lengths of size (bsz)
       tgt_in: a tensor of size (bsz, tgt_len)
  src_mask = src.ne(self.padding_id_src) # bsz, max_src_len
  # Forward encoder
  memory_bank, encoder_final_state = self.forward_encoder(src, src_lengths)
  # Forward decoder
  logits = self.forward_decoder(encoder_final_state, tgt_in, memory_bank,_
⇔src mask)
  return logits
def forward_decoder_incrementally(self, prev_decoder_states, tgt_in_onestep,
                                    memory_bank, src_mask,
                                    normalize=True):
   11 11 11
  Forward the decoder for a single step with token `tgt_in_onestep`.
  This function will be used both in `forward_decoder` and in beam search.
  Note that bsz can be greater than 1.
  Arguments:
      prev_decoder_states: a tuple (prev_decoder_state, prev_context).__
→ `prev_context`
                             is `None` for the first step
       tqt_in_onestep: a tensor of size (bsz), tokens at one step
       memory_bank: a tensor of size (bsz, src_len, hidden_size), encoder⊔
\hookrightarrow outputs
                    at every position
       src_mask: a tensor of size (bsz, src_len): a boolean tensor, `False`_
\hookrightarrow where
                 src is padding (we disallow decoder to attend to those ⊔
\hookrightarrow places).
       normalize: use log\_softmax to normalize or not. Beam search needs to_{\sqcup}
\neg normalize,
                  while `forward_decoder` does not
  Returns:
       logits: log probabilities for `tgt_in_token` of size (bsz, V_tgt)
       decoder_states: (`decoder_state`, `context`) which will be used for the
                       next incremental update
       attn: normalized attention scores at this step (bsz, src_len)
  prev_decoder_state, prev_context = prev_decoder_states
  #TODO
  decoder_states = (decoder_state, context)
  if normalize:
```

```
logits = torch.log_softmax(logits, dim=-1)
  return logits, decoder_states, attn
def evaluate_ppl(self, iterator):
  """Returns the model's perplexity on a given dataset `iterator`."""
  # Switch to eval mode
  self.eval()
  total_loss = 0
  total words = 0
  for batch in iterator:
    # Input and target
    src = batch['src_ids']
                                         # bsz, max src len
    src_lengths = batch['src_lengths'] # bsz
    tgt_in = batch['tgt_ids'][:, :-1] # Remove <eos> for decode input_u
(y_0 = \langle bos \rangle, y_1, y_2)
    tgt_out = batch['tgt_ids'][:, 1:] # Remove <bos> as target
                                                                        \rightarrow y_2, y_3 = \langle eos \rangle
    # Forward to get logits
    logits = self.forward(src, src_lengths, tgt_in) # bsz, tgt_len, V_tgt
    # Compute cross entropy loss
    loss = self.loss_function(logits.reshape(-1, self.V_tgt), tgt_out.
→reshape(-1))
    total loss += loss.item()
    total_words += tgt_out.ne(self.padding_id_tgt).float().sum().item()
  return math.exp(total_loss/total_words)
def train_all(self, train_iter, val_iter, epochs=10, learning_rate=0.001):
  """Train the model."""
  # Switch the module to training mode
  self.train()
  # Use Adam to optimize the parameters
  optim = torch.optim.Adam(self.parameters(), lr=learning_rate)
  best_validation_ppl = float('inf')
  best_model = None
  # Run the optimization for multiple epochs
  for epoch in range(epochs):
    total_words = 0
    total loss = 0.0
    for batch in tqdm(train_iter):
      # Zero the parameter gradients
      self.zero_grad()
      # Input and target
      tgt = batch['tgt_ids']
                                           # bsz, max_tgt_len
      src = batch['src ids']
                                           # bsz, max src len
      src_lengths = batch['src_lengths'] # bsz
      tgt_in = tgt[:, :-1].contiguous() # Remove <eos> for decode input_
(y_0 = \langle bos \rangle, y_1, y_2)
```

```
tgt_out = tgt[:, 1:].contiguous() # Remove <bos> as target
                                                                            (y_1, \dots
y_2, y_3 = \langle eos \rangle
      bsz = tgt.size(0)
      # Run forward pass and compute loss along the way.
      logits = self.forward(src, src_lengths, tgt_in)
      loss = self.loss function(logits.view(-1, self.V tgt), tgt out.view(-1))
      # Training stats
      num tgt words = tgt out.ne(self.padding id tgt).float().sum().item()
      total_words += num_tgt_words
      total_loss += loss.item()
      # Perform backpropagation
      loss.div(bsz).backward()
      optim.step()
    # Evaluate and track improvements on the validation dataset
    validation_ppl = self.evaluate_ppl(val_iter)
    self.train()
    if validation_ppl < best_validation_ppl:</pre>
      best_validation_ppl = validation_ppl
      self.best_model = copy.deepcopy(self.state_dict())
    epoch loss = total loss / total words
    print (f'Epoch: {epoch} Training Perplexity: {math.exp(epoch_loss):.4f} '
            f'Validation Perplexity: {validation_ppl:.4f}')
```

```
[]: EPOCHS = 2 # epochs, we highly recommend starting with a smaller number like 1
LEARNING_RATE = 2e-3 # learning rate

# Instantiate and train classifier
model = AttnEncoderDecoder(hf_src_tokenizer, hf_tgt_tokenizer,
hidden_size = 64,
layers = 3,
).to(device)

model.train_all(train_iter, val_iter, epochs=EPOCHS,
learning_rate=LEARNING_RATE)
model.load_state_dict(model.best_model)
```

Since the task we consider here is very simple, we should expect a perplexity very close to 1.

```
[]: # Evaluate model performance, the expected value should be < 1.05
print (f'Test perplexity: {model.evaluate_ppl(test_iter):.3f}')

[]: grader.check("encoder_decoder_ppl")</pre>
```

1.2 Beam search decoding

We can reuse most of our beam search code in lab 4-4 here: we only need to modify the code a bit to pass in memory_bank and src_mask. For reference here is the same pseudo-code used in lab 4-4, where we want to decode a single example x of maximum length max_T using a beam size of K.

```
def beam_search(x, K, max_T):
2.
        finished = []
                             # for storing completed hypotheses
        # Initialize the beam
        beams = [Beam(hyp=(bos), score=0)] # initial hypothesis: bos, initial score: 0
3.
4.
        for t in [1..max_T] # main body of search over time steps
5.
            hypotheses = []
            # Expand each beam by all possible tokens y_{t+1}
6.
            for beam in beams:
                y_{1:t}, score = beam.hyp, beam.score
7.
8.
                for y_{t+1} in V:
9.
                    y_{1:t+1} = y_{1:t} + [y_{t+1}]
                    new_score = score + log P(y_{t+1} | y_{1:t}, x)
10.
11.
                    hypotheses.append(Beam(hyp=y_{1:t+1}, score=new_score))
            # Find K best next beams
12.
            beams = sorted(hypotheses, key=lambda beam: -beam.score)[:K]
            # Set aside finished beams (those that end in <eos>)
            for beam in beams:
13.
14.
                y_{t+1} = beam.hyp[-1]
15.
                if y_{t+1} == eos:
16.
                    finished.append(beam)
                    beams.remove(beam)
17.
            # Break the loop if everything is finished
18.
            if len(beams) == 0:
19.
                break
        return sorted(finished, key=lambda beam: -beam.score)[0] # return the best finished h
20.
```

Implement function beam_search in the code below. In addition to the predicted target sequence, this function also returns a list of attentions all_attns.

```
[]: # max target length
MAX_T = 15
class Beam():
    """
    Helper class for storing a hypothesis, its score and its decoder hidden state.
    """
    def __init__(self, decoder_state, tokens, score):
        self.decoder_state = decoder_state
        self.tokens = tokens
```

```
self.score = score
class BeamSearcher():
  Main class for beam search.
  def __init__(self, model):
   self.model = model
    self.bos id = model.bos id
    self.eos_id = model.eos_id
    self.padding_id_src = model.padding_id_src
    self.V = model.V_tgt
  def beam_search(self, src, src_lengths, K, max_T=MAX_T):
    Performs beam search decoding.
    Arguments:
        src: src batch of size (1, max_src_len)
        src_lengths: src lengths of size (1)
        K: beam size
        max_T: max possible target length considered
    Returns:
        a list of token ids and a list of attentions
    finished = \Pi
    all attns = []
    # Initialize the beam
    self.model.eval()
    #TODO - fill in `memory_bank`, `encoder_final_state`, and `init_beam` below
    memory_bank = ...
    encoder_final_state = ...
    init_beam = ...
    beams = [init_beam]
    with torch.no_grad():
      for t in range(max_T): # main body of search over time steps
        # Expand each beam by all possible tokens y \{t+1\}
        all_total_scores = []
        for beam in beams:
          y_1_to_t, score, decoder_state = beam.tokens, beam.score, beam.
 →decoder state
          y_t = y_1_{to_t[-1]}
          #TODO - finish the code below
          # \mathit{Hint}: you might want to use `model.forward_decoder_incrementally`\sqcup
 ⇒with `normalize=True`
```

```
src_mask = src.ne(self.padding_id_src)
         logits = ...
         decoder_state = ...
         attn = ...
         total_scores = ...
         all_total_scores.append(total_scores)
         all_attns.append(attn) # keep attentions for visualization
         beam.decoder_state = decoder_state # update decoder state in the beam
      all_total_scores = torch.stack(all_total_scores) # (K, V) when t>0, (1, )
\hookrightarrow V) when t=0
       # Find K best next beams
       # The code below has the same functionality as lines 6-12, but is more
\rightarrow efficient
      all_scores_flattened = all_total_scores.view(-1) # K*V when t>0, 1*V_{\sqcup}
\rightarrowwhen t=0
      topk_scores, topk_ids = all_scores_flattened.topk(K, 0)
      beam_ids = topk_ids.div(self.V, rounding_mode='floor')
      next_tokens = topk_ids - beam_ids * self.V
      new_beams = []
      for k in range(K):
         beam_id = beam_ids[k]
                                      # which beam it comes from
         y t plus 1 = next tokens[k] # which y \{t+1\}
         score = topk_scores[k]
        beam = beams[beam id]
         decoder_state = beam.decoder_state
         y_1_{to} = beam.tokens
         #TODO
         new beam = ...
         new_beams.append(new_beam)
      beams = new_beams
       # Set aside completed beams
       # TODO - move completed beams to `finished` (and remove them from
→ `beams`)
       # Break the loop if everything is completed
      if len(beams) == 0:
           break
   # Return the best hypothesis
  if len(finished) > 0:
    finished = sorted(finished, key=lambda beam: -beam.score)
    return [token.item() for token in finished[0].tokens], all_attns
  else: # when nothing is finished, return an unfinished hypothesis
    return [token.item() for token in beams[0].tokens], all_attns
```

```
[]: grader.check("beam_search")
```

Now we can use beam search decoding to predict the outputs for the test set inputs using the trained model. You should expect an accuracy close to 100%.

```
[]: DEBUG_FIRST = 10 # set to 0 to disable printing predictions
     K = 1
                       # beam size 1
     correct = 0
     total = 0
     # create beam searcher
     beam_searcher = BeamSearcher(model)
     for index, batch in enumerate(test_iter, start=1):
       # Input and output
       src = batch['src ids']
       src_lengths = batch['src_lengths']
       # Predict
      prediction, _ = beam_searcher.beam_search(src, src_lengths, K)
       # Convert to string
       prediction = hf_tgt_tokenizer.decode(prediction, skip_special_tokens=True)
      ground_truth = hf_tgt_tokenizer.decode(batch['tgt_ids'][0],__
      ⇔skip_special_tokens=True)
      if DEBUG_FIRST > index:
         src = hf_src_tokenizer.decode(src[0], skip_special_tokens=True)
         print (f'Source: {src}')
         print (f'Prediction:
                              {prediction}')
         print (f'Ground truth: {ground_truth}')
       if ground_truth == prediction:
         correct += 1
       total += 1
     print (f'Accuracy: {correct/total:.2f}')
```

2 Visualizing attention

We can visualize how each query distributes its attention scores over each source word.

```
[]: K = 1 # this code only works for beam size 1

# Create beam searcher
beam_searcher = BeamSearcher(model)
batch = next(iter(test_iter))
# Input and output
src = batch['src_ids']
src_lengths = batch['src_lengths']
```

```
# Predict and get attentions
prediction, all_attns = beam_searcher.beam_search(src, src_lengths, K)
all_attns = torch.stack(all_attns, 0)
# Convert to string
prediction = hf_tgt_tokenizer.decode(prediction, skip_special_tokens=True)
ground_truth = hf_tgt_tokenizer.decode(batch['tgt_ids'][0],__
 ⇔skip_special_tokens=True)
src = hf_src_tokenizer.decode(src[0], skip_special_tokens=True)
print (f'Source: {src}')
print (f'Prediction:
                       {prediction}')
print (f'Ground truth: {ground_truth}')
# Plot
fig, ax = plt.subplots(figsize=(8, 6))
ax.imshow(all_attns[:,0,:].detach().cpu())
ax.set_yticks(list(range(1+len(prediction.split()))));
ax.set_yticklabels(prediction.split() + ['eos']);
ax.set_xticks(list(range(len(src.split()))));
ax.set_xticklabels(src.split());
# Uncomment the line below if the plot does not show up
# Make sure to comment that before submitting to gradescope
# since there would be some autograder issues with plt.show()
#plt.show()
```

Do these attentions make sense? Do you see how the attention mechanism solves the bottleneck problem in vanilla seq2seq?

3 The transformer architecture

In RNN-based neural encoder-decoder models, we used recurrence to model the dependencies among words. For example, by running a unidirectional RNN from y_1 to y_t , we can consider the past history when predicting y_{t+1} . However, running an RNN over a sequence is a serial process: we need to wait for it to finish running from y_1 to y_t before being able to compute the outputs at y_{t+1} . This serial process cannot be parallelized on GPUs along the sequence length dimension: even during training where all y_t 's are available, we cannot compute the logits for y_t and the logits for y_{t+1} in parallel.

The attention mechanism provides an alternative, and most importantly, parallelizable solution. The transformer model completely gets rid of recurrence and only uses attention to model the dependencies among words. For example, we can use attention to incorporate the representations from y_1 to y_t when predicting y_{t+1} , simply by attending to their word embeddings. This is called decoder self-attention.

Question: By getting rid of recurrence and only using decoder self-attention, can we compute the logits for any two different words y_{t_1} and y_{t_2} in parallel at training time (only consider decoder for now)? Why?

Type your answer here, replacing this text.

Similarly, at the encoder side, for each word x_i , we let it attend to the embeddings of x_1, \ldots, x_S , to model the context in which x_i appears. This is called *encoder self-attention*. It is different from decoder self-attention in that here every word attends to all words, but at the decoder side, every word can only attend to the previous words (since the prediction of word y_t cannot use the information from any $y_{>t}$).

To incorporate source-side information at the decoder side, at each time step, we let the decoder attend to the top-layer encoder outputs, as we did in the RNN-based encoder-decoder model above. This is called *cross-attention*. Note that there's no initialization of decoder hidden state here, since we no longer use an RNN.

The process we describe above is only a single layer of attention. In practice, transformers stack multiple layers of attention and feedforward layers, using the outputs from the layer below as the inputs to the layer above, as shown in the illustration below.

In the above illustration, due to space limits, we omitted the details of encoder self-attention and decoder self-attention, and we describe it here, using encoder-self-attention at layer 0 as an example. First, we use three linear projections to project each hidden state $h_{0,i}$ to a query vector $q_{0,i}$, a key vector $k_{0,i}$, and a value vector $v_{0,i}$. Then at each position i, we use q_i as the query, and $\{(k_{0,j}, v_{0,j}) : j \in \{1, \dots, S\}\}$ as keys/values to produce a context vector $c_{0,i}$. Note that the keys/values are the same for different positions, and the only difference is that a different query vector is used for each position.

A clear difference between the transformer architecture and the RNN-based encoder decoder architecture is that there are no horizontal arrows in the transformer model: transformers only use position-wise operations and attention operations. The dependencies among words are **only introduced by the attention operations**, while the other operations such as feedforward, non-linearity, and normalization are position-wise, that is, they do not depend on other positions, and can thus be performed in parallel.

Question: In the above transformer model, if we shuffle the input words x_1, \dots, x_4 , would we get a different distribution over y? Why or why not?

Type your answer here, replacing this text.

Since the transformer model itself doesn't have any sense of position or order, we encode the position of each word in the sentence, and add it to the word embedding as part of the input representation, as illustrated below.

The illustrations above also omitted residual connections, which add the inputs to certain operations (such as attention and feedforward) to the outputs. More details can be found in the code below.

3.1 Causal attention mask

To efficiently train the transformer model, we want to batch the attention operations together such that they can be fully parallelized along the sequence length dimension. (The non-attention operations are position-wise so they are trivally parallelizable.) This is quite straightforward for encoder self-attention and decoder-encoder cross-attention given our batched implementation of the attention function. However, things are a bit trickier for the decoder: each word y_t attends

to t-1 previous words y_1, \dots, y_{t-1} , which means each word y_t has a different set of key-value pairs. Is it possible to batch them together?

The solution is to use attention masks. For every word y_t , we give it all key-value pairs at y_1, \ldots, y_T , and we disallow attending to future words $y_t, y_{t+1}, \ldots, y_T$ through an attention mask. (Recall that the attention function takes a mask argument.) We usually call this attention mask a causal attention mask, as it prevents the leakage of information from the future into the past. Since every y_t has the same set of (key, value) pairs, we can batch them and compute the context vectors using a single call to the function attention.

What should such a mask be? Implement the causal mask function below to generate this mask.

Hint: you might find torch.triu useful.

```
[]: #TODO - implement this function, which returns a causal attention mask

def causal_mask(T):
    """
    Generate a causal mask.
    Arguments:
        T: the length of target sequence
    Returns:
        mask: a T x T tensor, where `mask[i, j]` should be `True`
        if y_i can attend to y_{j-1} (there's a "-1" since the first
        token in decoder input is <bos>) and `False` if y_i cannot
        attend to y_{j-1}
    """
    mask = ...
    return mask.to(device)
```

```
[]: grader.check("causal_attention_mask")
```

We can visualize the attention mask and manually check if it's what we expected.

```
[]: fig, ax = plt.subplots(figsize=(8, 6))

T = 7
mask = causal_mask(T)
ax.imshow(mask.cpu())

# Uncomment the line below if the plot does not show up
# Make sure to comment that before submitting to gradescope
# since there would be some autograder issues with `plt.show()`
#plt.show()
```

As we have emphasized multiple times, unlike RNN-based encoder-decoders, transformer encoder/decoders are parallelizable in the sequence length dimension, even for the decoder: by using causal masks, all positions (at the same layer) can be computed all at once (once the lower layer has been computed). The parallelizability of transformers is the key to its success, since it allows for training it on vast amounts of data.

Now we are ready to complete the implementation of the transformer model. The code is structured as a set of classes: TransformerEncoderLayer*, TransformerEncoder, TransformDecoderLayer*, TransformDecoder, PositionalEmbedding, and TransformerEncoderDecoder*. We've provided almost all the necessary code. In particular, we provide code for all position-wise operations. Your job is only to implement the parts involving attention and to figure out the correct attention masks, which involves only the three classes marked above with a star.

Hint: Completing this transformer implementation should require very little code, just a few lines.

Hint: The causal mask is a 2-D matrix, but we want to add a batch dimension, and expand it to be of the desired size. For this purpose, you can use torch.repeat.

```
[]: #TODO - implement `forward_encoder` and `forward_decoder`.
     # `TransformerEncoderDecoder` inherits most functions from `AttnEncoderDecoder`
     class TransformerEncoderDecoder(AttnEncoderDecoder):
       def __init__(self, hf_src_tokenizer, hf_tgt_tokenizer, hidden_size=64,_
      ⇒layers=3):
         n n n
         Initializer. Creates network modules and loss function.
         Arguments:
             hf_src_tokenizer: hf src tokenizer
             hf_tgt_tokenizer: hf tgt tokenizer
             hidden_size: hidden layer size of both encoder and decoder
             layers: number of layers of both encoder and decoder
         super(AttnEncoderDecoder, self).__init__()
         self.hf_src_tokenizer = hf_src_tokenizer
         self.hf_tgt_tokenizer = hf_tgt_tokenizer
         # Keep the vocabulary sizes available
         self.V_src = len(self.hf_src_tokenizer)
         self.V_tgt = len(self.hf_tgt_tokenizer)
         # Get special word ids or tokens
         self.padding_id_src = self.hf_src_tokenizer.pad_token_id
         self.padding_id_tgt = self.hf_tgt_tokenizer.pad_token_id
         self.bos_id = self.hf_tgt_tokenizer.bos_token_id
         self.eos_id = self.hf_tgt_tokenizer.eos_token_id
         # Keep hyper-parameters available
         self.embedding_size = hidden_size
         self.hidden_size = hidden_size
         self.layers = layers
         # Create essential modules
         self.encoder = TransformerEncoder(self.V_src, hidden_size, layers)
         self.decoder = TransformerDecoder(self.V_tgt, hidden_size, layers)
```

```
# Final projection layer
   self.hidden2output = nn.Linear(hidden_size, self.V_tgt)
  # Create loss function
  self.loss_function = nn.CrossEntropyLoss(reduction='sum',
                                              ignore_index=self.padding_id_tgt)
def forward encoder(self, src, src lengths):
  Encodes source words `src`.
  Arguments:
       src: src batch of size (bsz, max_src_len)
       src_lengths: src lengths (bsz)
  Returns:
       memory bank: a tensor of size (bsz, src_len, hidden_size)
   # The reason we don't directly pass in src_mask as in `forward decoder` is_
\hookrightarrow to
   # enable us to reuse beam search implemented for RNN-based encoder-decoder
  src len = src.size(1)
  #TODO - compute `encoder self attn mask`
  encoder_self_attn_mask = ...
  memory_bank = self.encoder(src, encoder_self_attn_mask)
  return memory_bank, None
def forward_decoder(self, tgt_in, memory_bank, src_mask):
   11 11 11
  Decodes based on memory bank, and ground truth target words.
  Arguments:
       tgt_in: a tensor of size (bsz, tgt_len)
       memory_bank: a tensor of size (bsz, src_len, hidden_size), encoder_\_
\hookrightarrow outputs
                    at every position
       src_mask: a tensor of size (bsz, src_len) which is `False` for source_
\hookrightarrow paddings
  Returns:
       Logits of size (bsz, tgt_len, V_tgt) (before the softmax operation)
  tgt_len = tgt_in.size(1)
  bsz = tgt_in.size(0)
   #TODO - compute `cross_attn_mask` and `decoder_self_attn_mask`
  cross_attn_mask = ...
  decoder_self_attn_mask = ...
  outputs = self.decoder(tgt_in, memory_bank, cross_attn_mask,_
→decoder_self_attn_mask)
```

```
logits = self.hidden2output(outputs)
  return logits
def forward(self, src, src_lengths, tgt_in):
  Performs forward computation, returns logits.
  Arguments:
       src: src batch of size (bsz, max_src_len)
       src_lengths: src lengths of size (bsz)
       tgt_in: a tensor of size (bsz, tgt_len)
  src_mask = src.ne(self.padding_id_src) # bsz, max_src_len
  # Forward encoder
  memory_bank, _ = self.forward_encoder(src, src_lengths)
  # Forward decoder
  logits = self.forward_decoder(tgt_in, memory_bank, src_mask)
  return logits
def forward_decoder_incrementally(self, prev_decoder_states, tgt_in_onestep,
                                    memory_bank, src_mask, normalize=True):
   11 11 11
  Forward the decoder at `decoder_state` for a single step with token_
\rightarrow 'tgt_in_onestep'.
   This function will be used in beam search. Note that the implementation \Box
⇔here is
  very inefficient, since we do not cache any decoder state, but instead we∟
   cache previously generated tokens in `prev_decoder_states`, and do a fresh
   `forward decoder`.
  Arguments:
       prev decoder states: previous tqt words. None for the first step.
       tgt_in_onestep: a tensor of size (bsz), tokens at one step
      memory_bank: a tensor of size (bsz, src_len, hidden_size), src hidden⊔
\hookrightarrowstates
                    at every position
       src_mask: a tensor of size (bsz, src_len): a boolean tensor, `False`_
\hookrightarrow where
                 src is padding.
       normalize: use log softmax to normalize or not. Beam search needs to \sqcup
\neg normalize,
                  while `forward_decoder` does not
  Returns:
       logits: Log probabilities for `tgt_in_token` of size (bsz, V_tgt)
       decoder_states: we use tgt words up to now as states, a tensor of size⊔
\hookrightarrow (bsz, len)
```

```
None: to keep output format the same as AttnEncoderDecoder, such that \sqcup
⇔we can
             reuse beam search code
  11 11 11
  prev tgt in = prev decoder states # bsz, tqt len
  src len = memory bank.size(1)
  bsz = memory_bank.size(0)
  tgt_in_onestep = tgt_in_onestep.view(-1, 1) # bsz, 1
  if prev_tgt_in is not None:
    tgt_in = torch.cat((prev_tgt_in, tgt_in_onestep), 1) # bsz, tgt_len+1
  else:
    tgt_in = tgt_in_onestep
  tgt_len = tgt_in.size(1)
  logits = self.forward_decoder(tgt_in, memory_bank, src_mask)
  logits = logits[:, -1]
  if normalize:
    logits = torch.log_softmax(logits, dim=-1)
  decoder_states = tgt_in
  return logits, decoder states, None
```

```
[]: class TransformerEncoder(nn.Module):
       r"""TransformerEncoder is an embedding layer and a stack of N encoder layers.
       Arguments:
           hidden size: hidden size.
           layers: the number of encoder layers.
       n n n
       def __init__(self, vocab_size, hidden_size, layers):
         super().__init__()
         self.embed = PositionalEmbedding(vocab_size, hidden_size)
         encoder_layer = TransformerEncoderLayer(hidden_size)
         self.layers = _get_clones(encoder_layer, layers)
         self.norm = nn.LayerNorm(hidden_size)
       def forward(self, src, encoder_self_attn_mask):
         r"""Pass the input through the word embedding layer, followed by
         the encoder layers in turn.
         Arguments:
             src: src batch of size (bsz, max_src_len)
             encoder\_self\_attn\_mask: the mask for encoder self\_attention, it's of_{\sqcup}
      \hookrightarrow size
                                       (bsz, max_src_len, max_src_len)
         Returns:
             a tensor of size (bsz, max_src_len, hidden_size)
```

```
output = self.embed(src)
    for mod in self.layers:
      output = mod(output, encoder_self_attn mask=encoder_self_attn mask)
    output = self.norm(output)
    return output
class TransformerEncoderLayer(nn.Module):
  r"""TransformerEncoderLayer is made up of self-attn and feedforward network.
  Arguments:
      hidden size: hidden size.
  def __init__(self, hidden_size):
    super(TransformerEncoderLayer, self).__init__()
    self.hidden_size = hidden_size
    fwd_hidden_size = hidden_size * 4
    # Create modules
    self.linear1 = nn.Linear(hidden_size, fwd_hidden_size)
    self.linear2 = nn.Linear(fwd_hidden_size, hidden_size)
    self.norm1 = nn.LayerNorm(hidden size)
    self.norm2 = nn.LayerNorm(hidden_size)
    self.activation = nn.ReLU()
    # Attention related
    self.q proj = nn.Linear(hidden size, hidden size)
    self.k_proj = nn.Linear(hidden_size, hidden_size)
    self.v_proj = nn.Linear(hidden_size, hidden_size)
    self.context_proj = nn.Linear(hidden_size, hidden_size)
  def forward(self, src, encoder_self_attn_mask):
    r"""Pass the input through the encoder layer.
    Arguments:
        src: an input tensor of size (bsz, max_src_len, hidden_size).
        encoder\_self\_attn\_mask: attention mask of size (bsz, max\_src\_len, \sqcup
 \hookrightarrow max\_src\_len),
                                 it's `False` where the corresponding attention _
 \hookrightarrow is disabled
    Returns:
        a tensor of size (bsz, max_src_len, hidden_size).
    q = self.q_proj(src) / math.sqrt(self.hidden_size) # a trick needed to make_
 ⇔transformer work
   k = self.k_proj(src)
    v = self.v_proj(src)
```

```
#TODO - compute `context`
    context = ...
    src2 = self.context_proj(context)
    # Residual connection
    src = src + src2
    src = self.norm1(src)
    # Feedforward for each position
    src2 = self.linear2(self.activation(self.linear1(src)))
    src = src + src2
    src = self.norm2(src)
    return src
class TransformerDecoder(nn.Module):
  r"""TransformerDecoder is an embedding layer and a stack of N decoder layers.
  Arguments:
      hidden_size: hidden size.
      layers: the number of sub-encoder-layers in the encoder.
  def __init__(self, vocab_size, hidden_size, layers):
    super(TransformerDecoder, self).__init__()
    self.embed = PositionalEmbedding(vocab size, hidden size)
    decoder_layer = TransformerDecoderLayer(hidden_size)
    self.layers = get clones(decoder layer, layers)
    self.norm = nn.LayerNorm(hidden_size)
  def forward(self, tgt_in, memory, cross_attn_mask, decoder_self_attn_mask):
    r"""Pass the inputs (and mask) through the word embedding layer, followed by
    the decoder layer in turn.
    Arguments:
        tqt_in: tqt batch of size (bsz, max_tqt_len)
        memory: the outputs of the encoder (bsz, max src len, hidden size)
        cross_attn_mask: attention mask of size (bsz, max_tqt_len,__
 \neg max\_src\_len),
                          it's `False` where the cross-attention is disallowed.
        decoder_self_attn_mask: attention mask of size (bsz, max_tqt_len,_
 \hookrightarrow max\_tgt\_len),
                                 it's `False` where the self-attention is_{\sqcup}
 \hookrightarrow disallowed.
    Returns:
        a tensor of size (bsz, max_tgt_len, hidden_size)
    output = self.embed(tgt_in)
    for mod in self.layers:
      output = mod(output, memory, cross_attn_mask=cross_attn_mask, \
                   decoder_self_attn_mask=decoder_self_attn_mask)
```

```
output = self.norm(output)
    return output
class TransformerDecoderLayer(nn.Module):
  r"""TransformerDecoderLayer is made up of self-attn, cross-attn, and
  feedforward network.
  Arguments:
      hidden size: hidden size.
  def __init__(self, hidden_size):
    super(TransformerDecoderLayer, self).__init__()
    self.hidden_size = hidden_size
    fwd_hidden_size = hidden_size * 4
    # Create modules
    self.linear1 = nn.Linear(hidden_size, fwd_hidden_size)
    self.linear2 = nn.Linear(fwd_hidden_size, hidden_size)
    self.activation = nn.ReLU()
    self.norm1 = nn.LayerNorm(hidden_size)
    self.norm2 = nn.LayerNorm(hidden size)
    self.norm3 = nn.LayerNorm(hidden_size)
    # Attention related
    self.q_proj_self = nn.Linear(hidden_size, hidden_size)
    self.k_proj_self = nn.Linear(hidden_size, hidden_size)
    self.v_proj_self = nn.Linear(hidden_size, hidden_size)
    self.context_proj_self = nn.Linear(hidden_size, hidden_size)
    self.q_proj_cross = nn.Linear(hidden_size, hidden_size)
    self.k_proj_cross = nn.Linear(hidden_size, hidden_size)
    self.v_proj_cross = nn.Linear(hidden_size, hidden_size)
    self.context_proj_cross = nn.Linear(hidden_size, hidden_size)
  def forward(self, tgt, memory, cross_attn_mask, decoder_self_attn_mask):
    r"""Pass the inputs (and mask) through the decoder layer.
    Arguments:
        tgt: an input tensor of size (bsz, max_tgt_len, hidden_size).
        memory: encoder outputs of size (bsz, max_src_len, hidden_size).
        cross_attn_mask: attention mask of size (bsz, max_tgt_len,_
 \neg max\_src\_len),
                         it's `False` where the cross-attention is disallowed.
        decoder_self_attn_mask: attention mask of size (bsz, max_tqt_len,_
 \hookrightarrow max_tgt_len),
```

```
it's `False` where the self-attention is_{\sqcup}
 \hookrightarrow disallowed.
    Returns:
        a tensor of size (bsz, max_tgt_len, hidden_size)
    # Self attention (decoder-side)
    q = self.q_proj_self(tgt) / math.sqrt(self.hidden_size)
    k = self.k_proj_self(tgt)
    v = self.v_proj_self(tgt)
    #TODO - compute `context`
    context = ...
    tgt2 = self.context_proj_self(context)
    tgt = tgt + tgt2
    tgt = self.norm1(tgt)
    # Cross attention (decoder attends to encoder)
    q = self.q_proj_cross(tgt) / math.sqrt(self.hidden_size)
    k = self.k_proj_cross(memory)
    v = self.v_proj_cross(memory)
    #TODO - compute `context`
    context = ...
    tgt2 = self.context proj cross(context)
    tgt = tgt + tgt2
    tgt = self.norm2(tgt)
    tgt2 = self.linear2(self.activation(self.linear1(tgt)))
    tgt = tgt + tgt2
    tgt = self.norm3(tgt)
    return tgt
class PositionalEmbedding(nn.Module):
  """Embeds a word both by its word id and by its position in the sentence."""
  def __init__(self, vocab_size, embedding_size, max_len=1024):
    super(PositionalEmbedding, self).__init__()
    self.embedding_size = embedding_size
    self.embed = nn.Embedding(vocab_size, embedding_size)
    pe = torch.zeros(max_len, embedding_size)
    position = torch.arange(0, max_len).unsqueeze(1)
    div_term = torch.exp(torch.arange(0, embedding_size, 2) *
                         -(math.log(10000.0) / embedding_size))
    pe[:, 0::2] = torch.sin(position * div_term)
    pe[:, 1::2] = torch.cos(position * div_term)
    pe = pe.unsqueeze(0) # 1, max_len, embedding_size
    self.register_buffer('pe', pe)
  def forward(self, batch):
    x = self.embed(batch) * math.sqrt(self.embedding_size) # type embedding
    # Add positional encoding to type embedding
```

```
x = x + self.pe[:, :x.size(1)].detach()
return x

def _get_clones(module, N):
    """Copies a module `N` times"""
return nn.ModuleList([copy.deepcopy(module) for i in range(N)])
```

You might notice that in these experiments training transformers doesn't appear to be faster than training RNNs. There are two reasons for that: first, we are not using GPUs; second, even if you use GPUs, the sequences here are too short to observe the benefits of parallelizing along the horizontal direction. In real datasets with long sentences, training transformers is much faster than training RNNs, so under the same computational budget, using transformers allows for training on much larger datasets. This is one of the primary reasons transformers dominate NLP research these days.

Question: We argued above that *training* transformers can be much faster than training RNNs. What about *generation* using transformers? Would there be any speed advantage of decoding (generation) using transformers compared to RNNs? Why or why not?

Type your answer here, replacing this text.

```
[]: # Evaluate model performance, the expected value should be < 1.5 print (f'Test perplexity: {model_transformer.evaluate_ppl(test_iter):.3f}')
```

```
[]: grader.check("transformer_ppl")
```

Now that we have a trained model, we can decode from it using our previously implemented beam search function. If the code below throws any errors, you might need to modify your beam search code such that it generalizes here.

```
[]: grader.check("transformer_beam_search")
```

```
[]: DEBUG_FIRST = 10 # set to False to disable printing predictions
     K = 1 \# beam size 1
     correct = 0
     total = 0
     # create beam searcher
     beam_searcher = BeamSearcher(model_transformer)
     for index, batch in enumerate(test_iter, start=1):
       # Input and output
       src = batch['src ids']
       src_lengths = batch['src_lengths']
       # Predict
      model.all_attns = []
      prediction, _ = beam_searcher.beam_search(src, src_lengths, K)
       # Convert to string
       prediction = hf_tgt_tokenizer.decode(prediction, skip_special_tokens=True)
       ground_truth = hf_tgt_tokenizer.decode(batch['tgt_ids'][0],__
      ⇔skip_special_tokens=True)
       if DEBUG_FIRST > index:
         src = hf_src_tokenizer.decode(src[0], skip_special_tokens=True)
         print (f'Source: {src}')
         print (f'Prediction:
                                {prediction}')
         print (f'Ground truth: {ground_truth}')
       if ground_truth == prediction:
         correct += 1
       total += 1
     print (f'Accuracy: {correct/total:.2f}')
```

Question: When we first introduced attention above, adding it to an RNN model, we noted that

The attention scores **a** lie on a *simplex* (meaning $a_i \ge 0$ and $\sum_i a_i = 1$), which lends it some interpretability: the closer a_i is to 1, the more "relevant" a key k_i (and hence its value v_i) is to the given query. We will observe this later in the lab: When we are about to predict the target word "3", a_i is close to 1 for the source word $x_i =$ "three".

Can we interpret the attentions in a multi-layer transformer similarly? If so, what would you expect the attention scores to correspond to? If not, explain why.

Type your answer here, replacing this text.

You might have noticed that the transformer model underperforms the RNN-based encoder-decoder on this particular task. This might be due to several reasons:

- Transformers tend to be data hungry, sometimes requiring billions of words to train.
- The transformer formulation presented in this lab is not in its full form: for instance, instead of only doing attention once at each position for each layer, researchers usually use multiple attention operations in the hope of capturing different aspects of "relevance", which is called

- "multi-headed attention". For example, one attention head might be focusing on pronoun resolution, while the other might be looking for similar contexts before.
- Transformers are usually sensitive to hyper-parameters and require heavy tuning. For example, while we used a fixed learning rate, researchers usually use a customized learning rate scheduler which first warms up the learning rate, and then gradually decreases it. If you are interested, more details can be found in the original paper.

In real-world applications, many state-of-the-art NLP approaches are based on transformers, such as the fake news generator used by GROVER that you've seen in the Embedded EthiCS class. For further readings if you are interested, we recommend BERT and GPT-3.

4 Lab debrief

Question: We're interested in any thoughts you have about this lab so that we can improve this lab for later years, and to inform later labs for this year. Please list any issues that arose or comments you have to improve the lab. Useful things to comment on might include the following, but you're not restricted to these:

- Was the lab too long or too short?
- Were the readings appropriate for the lab?
- Was it clear (at least after you completed the lab) what the points of the exercises were?
- Are there additions or changes you think would make the lab better?

Type your answer here, replacing this text.

End of Lab 4-5

To double-check your work, the cell below will rerun all of the autograder tests.

[]: grader.check_all()